

Problem set 42 - Intersection over Union (IoU) for object detection

For two given sets of boxes, your implementation should return IoUs for each pair of boxes (one box from each set).

Problem 42.1 - Optimize the worst case performance

The functionality should be implemented using a function with the following signature:

```
def iou(boxes_1: np.ndarray, boxes_2: np.ndarray) -> np.ndarray:  
    ...
```

Expect the boxes to be of shape $(n, 4)$ where n stands for the number of boxes in the given array, and 4 stands for (y_1, x_1, y_2, x_2) ((y_1, x_1) forms a topmost corner of the box, and (x_2, y_2) forms bottommost corner).

You can safely assume well defined boxes ($y_1 < y_2$ and $x_1 < x_2$).

Return a matrix of shape (n, m) where n and m stand for the number of boxes in `boxes_1` and `boxes_2` respectively. Each element of the returned matrix ((i, j)) should represent an IoU value for one of the pairs of boxes (`boxes_1[i]` and `boxes_2[j]` where $0 \leq i < n$ and $0 \leq j < m$).

Both inputs and outputs are of type `float32`.

Given $n, m < 2000$ your implementation should perform in under a second in [Google Colab](#) in a worst case scenario

Problem 42.2 - Matrix sparsity

Output matrix defined in 42.1 can be pretty [sparse](#) for some subspace of inputs.

Variables `H` and `W` are introduced such that $0 < y_1 < y_2 < H$ and $0 < x_1 < x_2 < W$ holds for each box.

Assume both sets of boxes come from the same distribution.

Using (x_1, y_1, h, w) notation of the box, where $h := y_2 - y_1$ and $w := x_2 - x_1$, assume h and w follow log-normal distribution.

Assume `x1` and `y1` follow uniform distribution along their domain.

Express sparsity (or density) of the output matrix in terms of `n`, `m` and the expectation of the terms `h / H` and `w / W`.

Comment the case where `h << H` and `w << W`.

Problem 42.3 - Optimize the average case performance

Assuming that expectation of the terms `h / H` and `w / W` is less than `1e-2`, with the maximal `n` and `m` from the Problem 42.1, reimplement the algorithm to optimize the average case performance. Feel free to use other output format as well.

General requirements

- use `>=Python3.6`
- only `numpy` can be installed as a dep
- type your code where suitable (`typing` package)
- implement a solution as a module (a single file)
- for coding problems send us the solution using [gist](#) or similar public snippet sharing service
- explicit control flow and loop constructs (`for`, `while`, `if`, ...) should be avoided
- code has to be both clean and performant
- for theoretic problems write your solution using LaTeX and send us the PDF